

Stream Processing with Apache Kafka & .NET

Niels Berglund - Architect Lead Derivco, Data Platform MVP

niels.it.berglund@gmail.com

<https://nielsberglund.com>

<https://twitter.com/nielsberglund>



Agenda

- Kafka 101
- Kafka & .NET
- Stream processing
- ksqlDB

Data

*Civilization has always run on data **

*Today, every aspect of human life is fueled by data **

*Data has become a renewable resource **

*Even when trade slows, data continuous to grow at a steady pace **

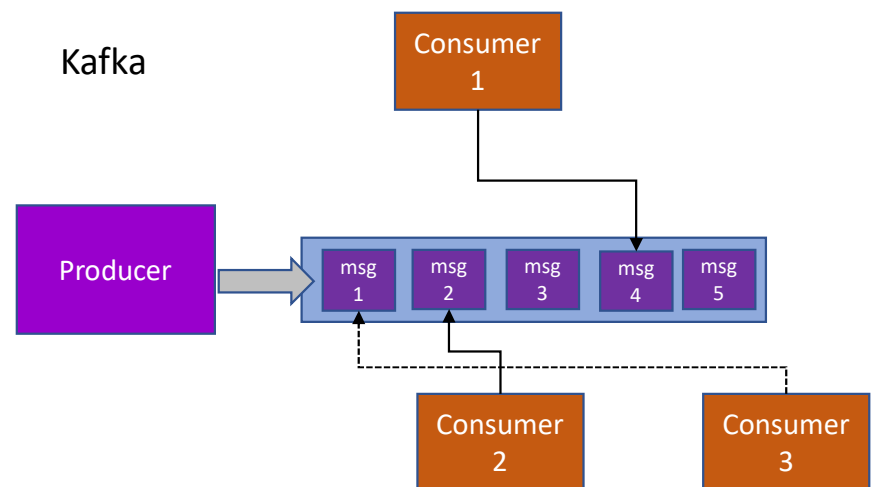

- 2016 - 16 Zettabytes of data
- 2025 - 165 Zettabytes of data

* Brad Smith: Tools and Weapons - The Promise and the Peril of the Digital Age

Kafka 101 - What is Kafka?

- Started at LinkedIn to replace their messaging systems.
 - Open sourced 2014
 - Confluent the commercial arm of Kafka
- Distributed streaming platform.
 - Publish and consume streams of records (events).
 - Persist published records.
 - Process streams of records in real-time.
- Commit log.
 - Logs are trustworthy.
 - Distributed for scale and resilience.

Kafka vs. Message Queues





Kafka 101 - Key Concepts

- Cluster
- Broker
- Topic
- Partition
- Offset
- Consumer groups

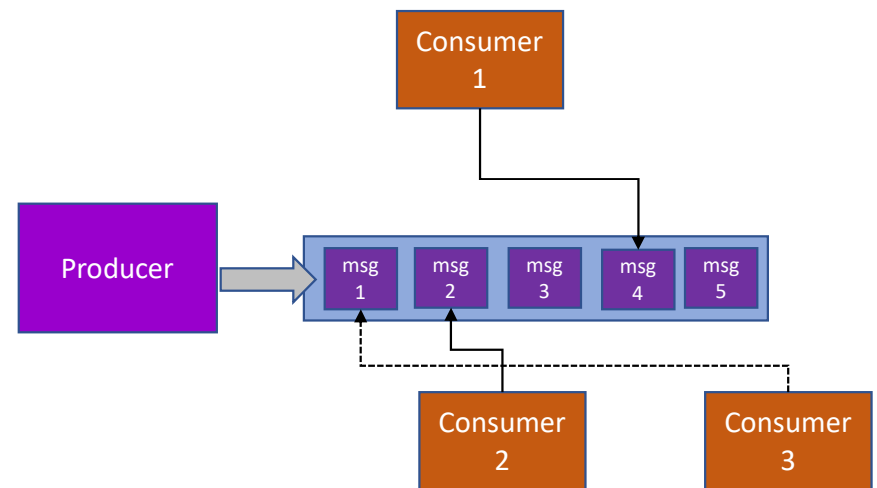
Kafka 101 - Topics & Partitions

- Topics
 - Name/definition of a stream of msgs/records/events, (wager, login, order, etc.).
 - Schema (JSON, Avro, CSV, Protobuf).
 - A message passed in to a topic has: key, value, timestamp.
 - Properties of a topic: retention period, compression type, number of partitions, replication factor.
- Partitions
 - A topic has partitions, (1 - n).
 - A partition is the actual log.
 - Partition provides scalability, parallelism, order.
 - Publisher controls the partitioning, (the key in the message).




Kafka 101 - Offsets

- Offset
 - Every record within a partition assigned a sequential ID.
 - Offset unique within the partition.
 - Consumer use offsets to indicate where to start/continue read from.
 - Provides a way to replay data.



Kafka 101 - Consumer Groups

- Consumer groups
 - Group of related consumers, reading from one or more partitions in one or more topics.
 - Only one reader from one consumer group allowed to read from a partition.
 - Multiple consumer groups can read from the same partition.



Kafka & .NET

- Kafka is Scala/Java based.
 - native Java client.
- Other languages uses librdkafka.dll.
 - librdkafka high performance C implementation of the Apache Kafka client.
- Latest version of .NET client on par with Java.
- Install the client from NuGet.

```
$ mkdir test  
$ cd test  
$ dotnet new console -f NET5.0  
# in VS Code terminal  
$ dotnet add package Confluent.Kafka --version 1.7.0
```

Kafka 101 - Produce

```
using System;
using System.Threading.Tasks;
using Confluent.Kafka;

class Program
{
    public static async Task Main(string[] args)
    {
        var config = new ProducerConfig { BootstrapServers = "localhost:9092" };
        using (var p = new ProducerBuilder<string, string>(config).Build())
        {
            try
            {
                var dr = await p.ProduceAsync("test-topic", new Message<Null, string> { Value="test" });
                Console.WriteLine($"Delivered '{dr.Value}' to '{dr.TopicPartitionOffset}'");
            }
            catch (ProduceException<Null, string> e)
            {
                Console.WriteLine($"Delivery failed: {e.Error.Reason}");
            }
        }
    }
}
```

<https://nielsberglund.com>

Kafka 101 - Consume

```
using System;
using System.Threading;
using Confluent.Kafka;

class Program
{
    public static void Main(string[] args)
    {
        var conf = new ConsumerConfig
        {
            GroupId = "test-consumer-group",
            BootstrapServers = "localhost:9092",
            AutoOffsetReset = AutoOffsetReset.Earliest
        };

        using (var c = new ConsumerBuilder<Ignore, string>(conf).Build())
        {
            c.Subscribe("my-topic");

            try
            {
                while (true)
                {
                    try
                    {
                        //cancellationtoken cts created above somewhere
                        var cr = c.Consume(cts.Token);
                        Console.WriteLine($"Consumed message '{cr.Message.Value}' at: '{cr.TopicPartitionOffset}'");
                    }
                    catch (ConsumeException e)
                    {
                        Console.WriteLine($"Error occurred: {e.Error.Reason}");
                    }
                }
            }
            catch (OperationCanceledException)
            {
                c.Close();
            }
        }
    }
}
```


Stream Processing

- Stream processing is what we do when we want to handle events




Stream Processing - I

authorization_attempts



possible_fraud




- Possible fraud = Authorization attempts > 3 within x timeframe

How to Implement Stream Processing

- Roll your own using Kafka consumer ☹️
 - how to handle state
- KStreams introduced 2016.
 - library for building streaming applications, specifically applications that transform input Kafka topics into output Kafka topics.
 - Scala/Java only ☹️
- KSQL introduced 2017
 - SQL Syntax. No Scala/Java! 😊

ksqlDB

- Evolution of KSQL, introduced 2019.
- Event streaming db.
- Allows you to build stream processing applications.
- Write real-time applications in SQL (no Java, Scala, etc.).
 - For complex functionality write User Defined Functions when needed.
- Lowers the bar for implementing streaming.



Lower the Bar for Streaming - I



ksqlDB

```
CREATE STREAM  
fraudulent_payments AS  
SELECT * FROM payments  
WHERE fraudProbability > 0.8;
```





```
object FraudFilteringApplication extends App {  
  
  val config = new java.util.Properties  
  config.put(StreamsConfig.APPLICATION_ID_CONFIG, "fraud-filtering")  
  config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "kafka-broker")  
  
  val builder: StreamsBuilder = new StreamsBuilder()  
  val fraudulentPayments: KStream[String, Payment] = builder  
    .stream[String, Payment]("payments-kafka-topic")  
    .filter((_, payment) => payment.fraudProbability > 0.8)  
  
  val streams: KafkaStreams = new KafkaStreams(builder.build(), config)  
  streams.start()  
}
```

Source: <https://www.slideshare.net/ConfluentInc/introduction-to-ksql-streaming-sql-for-apache-kafka>

<https://nielsberglund.com>

Lower the Bar for Streaming - II



ksqlDB - Streams and Tables

- Stream - unbounded sequence of structured data (events).
 - Immutable.
- Table - view of a stream or another table.
 - Mutable.
- Stream & table duality
 - Turn a stream into a table via DML manipulations: `SELECT`, `COUNT()`, `SUM()`, etc.
 - Turn a table into a stream by capturing inserts, updates, and deletes.

ksqlDB - Syntax

- Semantics similar to ANSI SQL.
- Statements:
 - CREATE ... (stream, table, and more).
 - Is combined with SELECT/INSERT.
 - Supports windowing functions, (SESSION, HOPPING, TUMBLING).
 - INSERT ...
 - SELECT (more later).
 - and more.
- Functions:
 - Scalar.
 - Aggregation.
 - Table.

ksqlDB - Push & Pull Queries

- Pull: returns the current value and terminates.

```
SELECT *  
FROM wagers  
WHERE ROWKEY = '12345';
```

- Push: returns a continuous stream of updates to a ksqlDB stream or table.

```
SELECT *  
FROM wagers  
EMIT CHANGES;
```

ksqlDB - Examples

```
CREATE STREAM fraudulent_payments  
AS  
SELECT * FROM payments  
WHERE fraudProbability > 0.8;
```

```
CREATE STREAM syslog_invalid_users  
AS  
SELECT host, message  
FROM syslog  
WHERE message LIKE '%Invalid user%';
```

```
CREATE TABLE possible_fraud  
AS  
SELECT card_number, COUNT(*)  
FROM authorization_attempts  
WINDOW TUMBLING (SIZE 5 SECONDS)  
GROUP BY card_number  
HAVING COUNT(*) > 3;
```

Summary

- Kafka distributed streaming platform.
- Topics, partitions, offsets, and consumer groups.
- Fully fledged .NET Client
- Stream processing handling the messages that arrives in Kafka
- ksqlDB lowers the barrier for entry to write real-time stream processing applications.

Thanks!

Questions?

Niels Berglund

niels.it.berglund@gmail.com

<https://nielsberglund.com>

<https://twitter.com/nielsberglund>